




Explore  
(/explore)




Tracks  
(/tracks)



My Courses  
(/mycourses)



Edpresso  
(/edpresso)



Refer a  
Friend  
(/refer-a-  
friend)



Create  
courses/grokking-  
system-design-  
view/3j6NnJrp

## Grokking the System Design Interview

(/collection/5668639101419520/56490502)

87% completed

Q Search Course

### Glossary of System Design Basics

- System Design Basics  
(/courses/grokking-the-system-design-interview/B892KY261z2)
- Key Characteristics of Distributed Systems  
(/courses/grokking-the-system-design-interview/YQWGjI ZZVz9)
- Load Balancing  
(/courses/grokking-the-system-design-interview/3jEwI04BL7Q)
- Caching  
(/courses/grokking-the-system-design-interview/3j6NnJrpp5p)
- Data Partitioning  
(/courses/grokking-the-system-design-interview/mEN8IJXV1LA)
- Indexes  
(/courses/grokking-the-system-design-interview/gxkVE8NEvXj)
- Proxies  
(/courses/grokking-the-system-design-interview/N8G9MvM4OR2)
- Redundancy and Replication  
(/courses/grokking-the-system-design-interview/xV1qvj6PKk)
- SQL vs. NoSQL  
(/courses/grokking-the-system-design-interview/YQIK1mDPgpK)
- CAP Theorem  
(/courses/grokking-the-system-design-interview/RMkqx1Egxqz)
- Consistent Hashing  
(/courses/grokking-the-system-design-interview/B81vnyp0GpY)
- Long-Polling vs WebSockets vs Server-Sent Events  
(/courses/grokking-the-system-design-interview/gx7wZzWn5Vj)

### Appendix

- Contact Us  
(/courses/grokking-the-system-design-interview/7nnxwPxΔI 7i)

## Caching

We'll cover the following

- Application server cache
- Content Distribution Network (CDN)
- Cache Invalidation
- Cache eviction policies

Load balancing helps you scale horizontally across an ever-increasing number of servers, but caching will enable you to make vastly better use of the resources you already have as well as making otherwise unattainable product requirements feasible. Caches take advantage of the locality of reference principle: recently requested data is likely to be requested again. They are used in almost every layer of computing: hardware, operating systems, web browsers, web applications, and more. A cache is like short-term memory: it has a limited amount of space, but is typically faster than the original data source and contains the most recently accessed items. Caches can exist at all levels in architecture, but are often found at the level nearest to the front end where they are implemented to return data quickly without taxing downstream levels.

### Application server cache

Placing a cache directly on a request layer node enables the local storage of response data. Each time a request is made to the service, the node will quickly return local cached data if it exists. If it is not in the cache, the requesting node will query the data from disk. The cache on one request layer node could also be located both in memory (which is very fast) and on the node's local disk (faster than going to network storage).

What happens when you expand this to many nodes? If the request layer is expanded to multiple nodes, it's still quite possible to have each node host its own cache. However, if your load balancer randomly distributes requests across the nodes, the same request will go to different nodes, thus increasing cache misses. Two choices for overcoming this hurdle are global caches and distributed caches.

### Content Distribution Network (CDN)

**Write-back cache:** Under this scheme, data is written to cache alone and completion is immediately confirmed to the client. The write to the permanent storage is done after specified intervals or under certain conditions. This results in low latency and high throughput for write-intensive applications, however, this speed comes with the risk of data loss in case of a crash or other adverse

PK

educative(/learn)

Explore(/explore)

Tracks(/tracks)

My Courses(/mycourses)

Edpresso(/edpresso)

Refer a Friend(/refer-a-friend)

Create Courses(/courses/grokking-the-system-design-view/3j6NnJrp)

Grokking the System Design Interview

(/collection/5668639101419520/56490502)

87% completed

Q Search Course

System Design Basics(/courses/grokking-the-system-design-interview/B892KY261z2)

Key Characteristics of Distributed Systems(/courses/grokking-the-system-design-interview/YQWGjIzzVz9)

Load Balancing(/courses/grokking-the-system-design-interview/3jEwI04BL7Q)

Caching(/courses/grokking-the-system-design-interview/3j6NnJrpp5p)

Data Partitioning(/courses/grokking-the-system-design-interview/mEN8IJXV1LA)

Indexes(/courses/grokking-the-system-design-interview/gxkVE8NEvXj)

Proxies(/courses/grokking-the-system-design-interview/N8G9MvM4OR2)

Redundancy and Replication(/courses/grokking-the-system-design-interview/xV1qvj6PKk)

SQL vs. NoSQL(/courses/grokking-the-system-design-interview/YQIK1mDPgpK)

CAP Theorem(/courses/grokking-the-system-design-interview/RMkqx1Egxqz)

Consistent Hashing(/courses/grokking-the-system-design-interview/B81vnyp0GpY)

Long-Polling vs WebSockets vs Server-Sent Events(/courses/grokking-the-system-design-interview/gx7wZzWn5Vj)

Appendix

Contact Us(/courses/grokking-the-system-design-interview/7nnxwPxAl 7il)

event because the only copy of the written data is in the cache.

Cache eviction policies

#

Following are some of the most common cache eviction policies:

1. First In First Out (FIFO): The cache evicts the first block accessed first without any regard to how often or how many times it was accessed before.

2. Last In First Out (LIFO): The cache evicts the block accessed most recently first without any regard to how often or how many times it was accessed before.

3. Least Recently Used (LRU): Discards the least recently used items first.

4. Most Recently Used (MRU): Discards, in contrast to LRU, the most recently used items first.

5. Least Frequently Used (LFU): Counts how often an item is needed. Those that are used least often are discarded first.

6. Random Replacement (RR): Randomly selects a candidate item and discards it to make space when necessary.

Following links have some good discussion about caching:

[1] Cache (https://en.wikipedia.org/wiki/Cache\_(computing))

[2] Introduction to architecting systems (https://lethain.com/introduction-to-architecting-systems-for-scale/)

← Back (/courses/grokking-the-system-design-interview/3jEwI04BL7Q)

Next → 

COMPLETED (/courses/grokking-the-system-design-interview/mEN8IJXV1LA)

Report an Issue

Ask a Question (https://discuss.educative.io/c/grokking-the-system-design-interview-design-gurus/glossary-of-system-design-basics-caching)